



1. GridGain In-Memory Accelerator For Apache Hadoop

GridGain's In-Memory Accelerator For Apache Hadoop is designed to deliver uncompromised performance for existing Apache Hadoop 2.* applications with zero code change as well as simplicity of installation and configuration across all the supported platforms.

2. Installation

GridGain distribution comes in a ZIP file that simply needs to be unzipped. The Accelerator requires Apache Hadoop of version 2.* to be already installed on the system either using Apache Bigtop packages or manually (manual installation just means that Apache Hadoop binary distribution must be unpacked somewhere on the system). In case of manual installation `HADOOP_HOME` environment variable must point to the installation directory of Apache Hadoop.

The Accelerator comes with command line setup tool `bin/setup-hadoop.sh` (`bin/setup-hadoop.bat` on Windows) which will guide you through all the needed setup steps (note that the setup tool will require write permissions to the Apache Hadoop installation directory).

Installation requirements:

1. Windows, Linux, or MacOS environment.
2. Java 7 or 8 (latest update is advisable).
3. Point `JAVA_HOME` environment variable to your JDK or JRE installation.
4. Apache Hadoop 2.* installed.
5. Point `HADOOP_HOME` environment variable to the installation directory of Apache Hadoop.
6. Run `bin/setup-hadoop.{sh|bat}` setup script and follow instructions.

2.1 Check GridGain Installation

After setup script successfully completed, you can execute the GridGain startup script. The following command will startup GridGain node with default configuration using Multicast node discovery.

```
bin/ggstart.{sh|bat}
```

If GridGain was installed successfully, the output from above commands should produce no exceptions or errors. If the Apache Hadoop installation was not found GridGain node will print respective warning. Note that you may see some other warnings during startup, but this is OK as they are meant to inform that certain functionality is turned on or off by default.

You can execute the above commands multiple times on the same machine and make sure that nodes discover each other. Here is an example of log printout when 2 nodes join topology:

```
... Topology snapshot [nodes=2, CPUs=8, hash=0xD551B245]
```

You can also start GridGain Management Console, called Visor, and observe started nodes show up on Visor Dashboard. To startup Visor, you should execute the following script:

```
bin/ggvisorcnd.{sh|bat} (Command-line, available in open source)
bin/ggvisorui.{sh|bat} (GUI mode, enterprise version only)
```

3. Configuration

To configure GridGain nodes you can change configuration files at `config` directory of GridGain installation. Those are conventional Spring files. Please refer to shipped configuration files and GridGain javadocs for more details.

3.1 Distributed File System Configuration

GridGain has its own distributed in-memory file system called GGFS. Hadoop jobs can use it instead of HDFS to achieve maximum performance and scalability. Setting up GGFS is much simpler than HDFS, it requires just few tweaks of GridGain node configuration and does not require starting any additional processes. Default configuration shipped with the

Accelerator contains one configured instance named "ggfs" which can be used as reference.

In common URI for GGFS looks like:

```
ggfs://ggfs_name@host_name
```

Where `ggfs_name` is GGFS instance name, `host_name` is any host running GridGain node with that GGFS instance configured. For more details please refer to GGFS documentation.

3.2 Apache Hadoop Client Configuration

To run Apache Hadoop jobs with GridGain cluster you need to configure `core-site.xml` and `mapred-site.xml` the same way as it is done in templates shipped with the Accelerator. You can find these templates at [docs/core-site.xml.gridgain](#) and [docs/mapred-site.xml.gridgain](#) respectively.

4. Running Apache Hadoop Job With GridGain's In-Memory Accelerator

When all the configuration is complete running Apache Hadoop job will be the same as with conventional Apache Hadoop distribution except that all GridGain nodes are equal and any of them can be treated as Job Tracker and DFS Name Node.

To run "Word Count" example you can load some text files to GGFS using standard Apache Hadoop tools:

```
$HADOOP_HOME/bin/hadoop fs -mkdir /input

$HADOOP_HOME/bin/hadoop fs -copyFromLocal /path/to/my/favorite/book.txt /input/my_book.txt
```

Run the job:

```
$HADOOP_HOME/bin/hadoop jar $HADOOP_HOME/share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar wordcount /input /output
```

Check results:

```
$HADOOP_HOME/bin/hadoop fs -ls /output

$HADOOP_HOME/bin/hadoop fs -cat /output/part-r-00000
```

A job can be ran on multiple nodes on localhost or in cluster environment the same way. The only changes needed to switch Apache Hadoop client to a cluster are to fix host in default DFS URI in `core-site.xml` and host in job tracker address in `mapred-site.xml`.

5. Management & Monitoring with Visor

GridGain comes with GUI and CLI (command) based DevOps Managements Consoles delivering advance set of management and monitoring capabilities. Visor GUI is based on a standalone Java application and CLI version is built on top of Scala REPL providing fully scriptable and customizable DevOps environment.

To start Visor GUI console (enterprise edition only), execute the following command:

```
`bin/ggvisorui.sh`
```

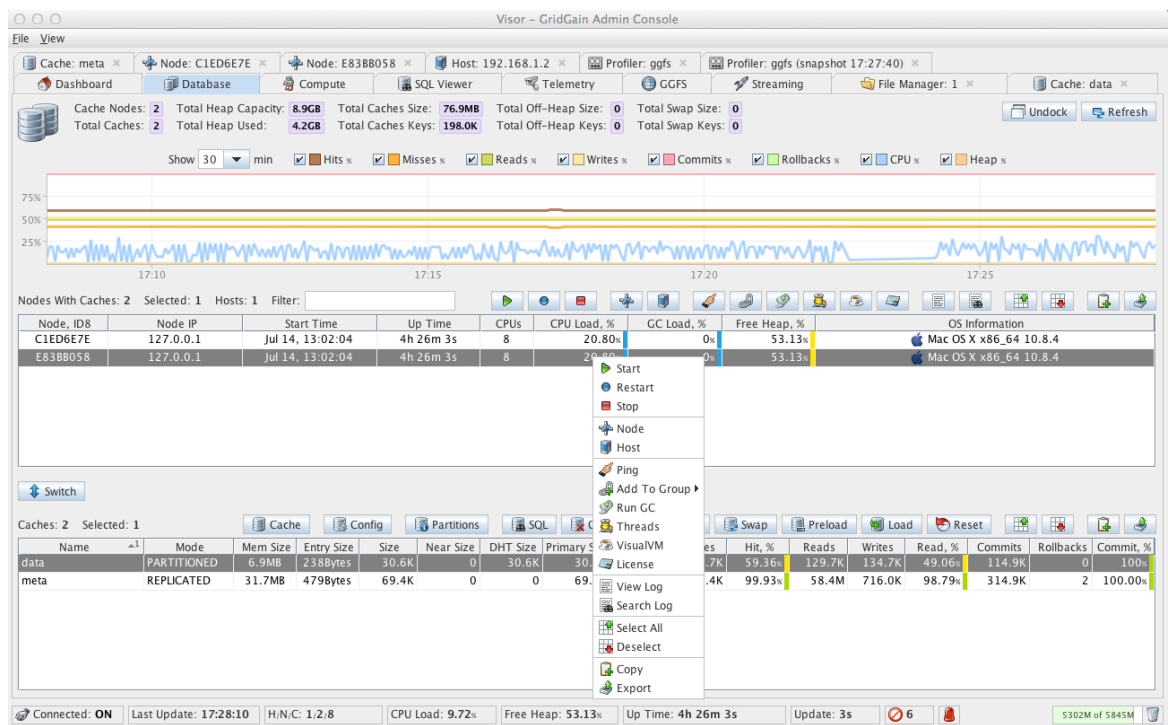
To start Visor in console mode you should execute the following command:

```
`bin/ggvisorcmd.sh`
```

On Windows, run the same commands with `.bat` extension.

NOTE: Visor GUI console has a much richer set of functionality over Visor command-based console. You should always prefer Visor GUI console whenever possible.

Here is an example of [Visor Data Grid Tab](#) which provides overall view on data grid.



Here is an example of [Visor Cache Tab](#) which provides view on individual cache.

